


Axis 1 — what is learned

classical / statistical → deep & purpose-built → LLM / foundation models

Scalable & interpretable methods

5 papers



§3.3

Platforms, pipelines & infrastructure


3 papers



§3.4

Applications & domains

5 papers



§3.5

Foundation models & LLMs

4 papers



§3.3 / §3.5

every paper sits in exactly one branch

Privacy, governance & evaluation

3 papers



§3.6

Empty cells (findings)

- no reinforcement-learning-at-scale paper
- no purely theoretical contribution
- no 2027-dated item

Axis 2 — what surrounds the learning

platform (B2) · domain (B3) · governance (B5)